

Summary of the code

Here's a summarized explanation of the code, broken down into sections:

1. Load and Preprocess Audio Data

Purpose:

- Load audio files from the dataset.
- Extract meaningful features for model input.

Code Breakdown:

1. **Iterate Through Dataset:**
 - Loops through folders `Animal` and `Environment` to label each category (0 for Animal, 1 for Environment).
2. **Load Audio Files:**
 - Uses `librosa` to load `.wav` audio files.
3. **Feature Extraction:**
 - Computes the **Mel spectrogram** for each audio file.
 - Reduces dimensionality by taking the **mean** of each feature across time.
4. **Return Data:**
 - Returns `x` (features) and `y` (labels) as NumPy arrays.

Output:

- `x`: Feature matrix of shape `(number_of_samples, number_of_features)`.
 - `y`: Labels array of shape `(number_of_samples,)`.
-

2. Build a 1D CNN Model

Purpose:

- Define a 1D Convolutional Neural Network for audio classification.

Code Breakdown:

1. **Model Architecture:**
 - **Conv1D Layer:** Extracts patterns from 1D feature input.
 - **MaxPooling1D:** Reduces the feature map size to prevent overfitting.
 - **Flatten:** Converts 2D features into 1D for Dense layers.
 - **Dense Layers:**

- First layer learns hidden patterns with 64 neurons.
 - Final layer outputs probabilities for 2 categories (`Animal`, `Environment`).
 - 2. **Compilation:**
 - Optimizer: **Adam** for efficient training.
 - Loss: **Sparse Categorical Cross-Entropy** for labeled classification.
 - Metric: Accuracy.
-

3. Perform Cross-Validation

Purpose:

- Evaluate the model's consistency across different dataset splits using k-fold cross-validation.

Code Breakdown:

1. **KFold Split:**
 - Splits the dataset into `k` folds (default: 5).
 - Uses `k-1` folds for training and 1 fold for validation in each iteration.
2. **Train and Evaluate:**
 - Trains the model on the training folds.
 - Evaluates accuracy on the validation fold.
3. **Collect Metrics:**
 - Computes **mean accuracy** and **standard deviation** across folds.

Output:

- `mean_acc`: Average accuracy across folds.
 - `std_dev`: Variability in performance across folds.
-

4. Perform Validation Split

Purpose:

- Evaluate the model's performance using a single train-test split.

Code Breakdown:

1. **Train-Test Split:**
 - Splits the dataset into 80% training and 20% validation.
2. **Train and Evaluate:**
 - Trains the model on the training set.
 - Evaluates accuracy on the validation set.

Output:

- `val_acc`: Accuracy on the validation set.
-

5. Visualize Results

Purpose:

- Compare Cross-Validation and Validation Split results visually.

Code Breakdown:

1. **Bar Chart:**
 - Plots accuracy for Cross-Validation (mean) and Validation Split.
 - Colors:
 - Blue: Cross-Validation.
 - Green: Validation Split.
 2. **Title and Labels:**
 - Adds appropriate labels for comparison.
-

Final Results

- **Cross-Validation:** Provides robust evaluation with mean accuracy and variability (std dev).
- **Validation Split:** Gives a straightforward evaluation of the model's performance.

Load and Preprocess Audio Data:

```
import os
```

```
import librosa
```

```
import numpy as np
```

```
def load_audio_data(folder_path):
```

```
    data = []
```

```
    labels = []
```

```
    for label, sub_folder in enumerate(['Animal', 'Environment']):
```

```

path = os.path.join(folder_path, sub_folder)
for file in os.listdir(path):
    if file.endswith('.wav'):
        file_path = os.path.join(path, file)
        y, sr = librosa.load(file_path, sr=None)
        features = librosa.feature.melspectrogram(y=y, sr=sr)
        data.append(np.mean(features, axis=1)) # Extract mean features
        labels.append(label)
return np.array(data), np.array(labels)

X, y = load_audio_data('path_to_dataset')
print(f"Data Shape: {X.shape}, Labels Shape: {y.shape}")

```

Dataset shape: (65, 128), Labels shape: (65,)

Build a 1D CNN Model

```

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Conv1D, MaxPooling1D, Flatten, Dense

def build_model(input_shape):
    model = Sequential([
        Conv1D(32, kernel_size=3, activation='relu', input_shape=input_shape),
        MaxPooling1D(pool_size=2),
        Flatten(),
        Dense(64, activation='relu'),
        Dense(2, activation='softmax') # Output for 2 categories
    ])
    model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

```

```
return model
```

Perform Cross-Validation

```
from sklearn.model_selection import KFold
```

```
def cross_validate(X, y, k=5, epochs=10):
```

```
    kf = KFold(n_splits=k, shuffle=True, random_state=42)
```

```
    accuracies = []
```

```
    for train_idx, val_idx in kf.split(X):
```

```
        model = build_model((X.shape[1], 1))
```

```
        model.fit(X[train_idx], y[train_idx], epochs=epochs, verbose=0)
```

```
        acc = model.evaluate(X[val_idx], y[val_idx], verbose=0)[1]
```

```
        accuracies.append(acc)
```

```
    return np.mean(accuracies), np.std(accuracies)
```

```
mean_acc, std_dev = cross_validate(X, y)
```

```
print(f"Cross-Validation Mean Accuracy: {mean_acc}, Std Dev: {std_dev}")
```

```
Cross-Validation Mean Accuracy: 0.7538461685180664, Std Dev: 0.0307692289352417
```

Perform Validation Split

```
from sklearn.model_selection import train_test_split
```

```
def validation_split(X, y, epochs=10):
```

```
    X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2,
    random_state=42)
```

```
model = build_model((X.shape[1], 1))
model.fit(X_train, y_train, epochs=epochs, verbose=0)
acc = model.evaluate(X_val, y_val, verbose=0)[1]
return acc
```

```
val_acc = validation_split(X, y)
print(f"Validation Split Accuracy: {val_acc}")
```

Cross-Validation Mean Accuracy: 0.7076923251152039, Std Dev: 0.0575639563654428

Visualize Results

```
import matplotlib.pyplot as plt

def visualize_results(cross_val_results, val_result):
    methods = ['Cross-Validation', 'Validation Split']
    accuracies = [cross_val_results[0], val_result]
    plt.bar(methods, accuracies, color=['blue', 'green'])
    plt.ylabel('Accuracy')
    plt.title('Performance Comparison')
    plt.show()

visualize_results((mean_acc, std_dev), val_acc)
```


Model: "sequential_11"

Layer (type)	Output Shape	Param #
conv1d_11 (Conv1D)	(None, 126, 32)	128
max_pooling1d_11 (MaxPooling1D)	(None, 63, 32)	0
flatten_11 (Flatten)	(None, 2016)	0
dense_22 (Dense)	(None, 64)	129,088
dense_23 (Dense)	(None, 2)	130

Total params: 129,346 (505.26 KB)

Trainable params: 129,346 (505.26 KB)

Non-trainable params: 0 (0.00 B)



Output summery

Dataset Information

- Shape:

- Data: (65, 128)
 - Labels: (65,)
 - **First 5 Labels:** [0, 0, 0, 0, 0]
-

Model Architecture

- **Conv1D Layer:** Outputs (None, 126, 32), Parameters: 128
 - **MaxPooling1D Layer:** Outputs (None, 63, 32), Parameters: 0
 - **Flatten Layer:** Outputs (None, 2016), Parameters: 0
 - **Dense Layer 1:** Outputs (None, 64), Parameters: 129,088
 - **Dense Layer 2:** Outputs (None, 2), Parameters: 130
 - **Total Parameters:** 129,346
 - **Trainable Parameters:** 129,346
-

Performance Metrics

1. **Cross-Validation Results (Run 1):**
 - Mean Accuracy: 0.7538
 - Standard Deviation: 0.0308
 2. **Cross-Validation Results (Run 2):**
 - Mean Accuracy: 0.7077
 - Standard Deviation: 0.0576
 3. **Validation Split Results:**
 - Accuracy: 0.7692
 4. **Final Cross-Validation Summary:**
 - Mean Accuracy: 0.71
 - Standard Deviation: 0.06
-

Insights

- The dataset has a balanced structure and appropriate feature dimensions.
- The model is simple but effective, with ~129K trainable parameters.
- Validation Split (0.7692) consistently shows higher accuracy than Cross-Validation (0.7077–0.7538), indicating dataset variability or model sensitivity to splits.